

H1 Report of CS489 Project 1

H2 0. Introduction

In this project, we are required to build the GridWorld environment and implement iterative policy evaluation method and policy iteration method. These are two methods based on Dynamic Programming.

H2 1. GridWorld Environment

Literally, GridWorld is a grid. Each grid in it represents a state. There're terminal states and non-terminal states. In non-terminal states, we can move one grid to north, east, south and west.

0	1	2	3	4	5
6	7	8	9	10	11
12	13	14	15	16	17
18	19	20	21	22	23
24	25	26	27	28	29
30	31	32	33	34	35

Fig.1 GridWorld

We can build an MDP $\langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma \rangle$ for GridWorld, in which we choose an random uniform policy.

- \mathcal{S} is a finite set of states, here $\mathcal{S} = \{s_t | t \in 0, \dots, 35\}$, s_1 and s_{35} are terminal states
- $\mathcal{A} = \{n, e, s, w\}$ which represents north, east, south and west move
- \mathcal{P} is a state transition probability matrix where

$$\mathcal{P}_{ss'}^a = \mathbb{P}[S_{t+1} = s' | S_t = s, A_t = a] = 1.0, \forall s \in \mathcal{S} - \{s_1, s_{35}\}, s' \in \mathcal{S}, a \in \mathcal{A}$$

$$\mathcal{P}_{ss'}^a = \mathbb{P}[S_{t+1} = s' | S_t = s, A_t = a] = 0, \forall s \in \{s_1, s_{35}\}, s' \in \mathcal{S}, a \in \mathcal{A}$$

- \mathcal{R} is a reward function where

$$\mathcal{R}_s^a = \mathbb{E}[R_{t+1} | S_t = s, A_t = a] = -1.0, \forall s \in \mathcal{S} - \{s_1, s_{35}\}, a \in \mathcal{A}$$

$$\mathcal{R}_s^a = \mathbb{E}[R_{t+1} | S_t = s, A_t = a] = 0, \forall s \in \{s_1, s_{35}\}, a \in \mathcal{A}$$

- $\gamma = 1.0$ for it's a episodic task
- π is the policy where

$$\pi(a|s) = \mathbb{P}[A_t = a | S_t = s] = 0.25, \forall a \in \mathcal{A}, s \in \mathcal{S}$$

And we can implement the GridWorld MDP in Python. The definition of data structures are shown as follows:

```
1 class GridWorld:
2     def __init__(self, state, terminalState, gamma, threshold):
3         self.state = state #|S|
4         self.terminalState = terminalState
5         self.gamma = gamma #gamma
```

```

6     self.threshold = threshold #theta
7
8     self.value = []
9     self.optimalPolicy = []
10    self.policy = [] #pi[s][a]
11    self.gridSize = int(math.sqrt(self.state))
12    self.action = {'n': 0, 'e': 1, 's': 2, 'w': 3} #A
13    self.trans = [[0 for j in range(len(self.action))] for i in range(self.state)] #s' =
trans[s][a]
14    self.prob = [[[0.0 for k in range(self.state)] for j in range(len(self.action))]
for i in range(self.state)] #P[s][a][s']
15    self.reward = [[-1.0 for j in range(len(self.action))] for i in range(self.state)]
#E[R[s][a]]
16    self.map = [[[j + i * self.gridSize) for j in range(self.gridSize)] for i in
range(self.gridSize)] #for calculating trans[s][a] and filling in P[s][a][s']

```

We can calculate \mathcal{P} by invoking `self.__calcParam` and fill the probability 1.0 in the right place in `self.prob`. The policy `self.policy` will be initialized later.

H2 2. Iterative Policy Evaluation

H3 2.1 Implementation

The policy evaluation is easy to implement using the data structures we declared before. We need to iteratively calculate the Bellman equation until converge. Thus, we can use the following equation to update our state value function $v(s)$

$$v_{k+1}(s) = \sum_{a \in \mathcal{A}} \pi(a|s) (\mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^a v_k(s'))$$

Before the iterative update of $v(s)$, we need to initialize $v(s)$ arbitrarily except $v(\text{terminal}) = 0$. And here's my implementation of the iteration

```

1  def evaluation(self):
2      #Initialize
3      .....
4      k = 0
5      #Loop
6      while True:
7          delta = 0.0
8          for curState in range(self.state):
9              oldValue = self.value[curState]
10             newValue = 0.0
11             for a in self.action.keys(): #GridWorld Specified
12                 tmp = 0.0
13                 #the inner sum disappeared
14                 nextState = self.trans[curState][self.action[a]] #Only one element
because here's GridWorld
15                 tmp += self.prob[curState][self.action[a]][nextState] *
self.value[nextState] #prob must be 1.0
16                 newValue += self.policy[curState][self.action[a]] *
(self.reward[curState][self.action[a]] + (self.gamma * tmp))
17             self.value[curState] = newValue
18             delta = max(delta, math.fabs(oldValue - self.value[curState]))

```

```

19     k += 1
20     if delta < self.threshold:
21         break

```

As we can see, in the inner for-loop we are calculating $\mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^a v_k(s')$. However, due to the property of GridWorld, the sum over $s' \in \mathcal{S}$ has only one term. Thus, no another for-loop is needed.

H3 2.2 Result

The algorithm on our GridWorld converges in about 300 iterations. And here's the result.

```

Iterative Policy Evaluation:
-18.17  0.00   -29.22  -44.06  -51.56  -54.68
-32.34  -30.17  -39.60  -47.41  -51.93  -53.80
-44.68  -44.74  -47.58  -50.06  -50.96  -50.79
-52.97  -52.51  -51.95  -50.27  -47.05  -43.61
-57.71  -56.38  -53.44  -48.01  -39.38  -29.00
-59.79  -57.86  -53.42  -44.96  -29.45   0.00

```

Fig.2 Iterative Policy Evaluation

H2 3. Policy Iteration

H3 3.1 Implementation

The policy iteration based on Bellman optimality equation is used to improve a given policy. I tried the algorithm shown in the slide. However, the program fell into a infinite loop. I found it will never converge if I arbitrarily initialize the policy as a deterministic policy. It's because the initialization may generate a deterministic policy which never reaches the terminal states. I figured out two methods to tackle with this problem.

1. Remove the outer while-loop, which means evaluate only once
2. Use nondeterministic policy

According to method 1, I rewrote the evaluation parts and it worked. It returned a optimal deterministic policy. However, I found there are some states which have more than one paths to the terminal states. It reminded me of method 2 and I modified the given algorithm.

```

1  def policyIteration(self):
2      #Initialize
3      .....
4      self.optimalPolicy = [[0, 1, 2, 3] for i in range(self.state)]
5      .....
6      numPass = 0
7      while True:
8          #Loop
9          .....
10
11         #Policy Improvement
12         policyStable = True
13         for curState in range(self.state):

```

```

14     oldAction = self.optimalPolicy[curState]
15     newAction = []
16     tmpValue = []
17     maxValue = -1.0e9
18     for a in self.action.keys():
19         nextState = self.trans[curState][self.action[a]] #Only one element
because here's GridWorld
20         tmp = self.prob[curState][self.action[a]][nextState] *
self.value[nextState] #prob must be 1.0
21         tmp = self.reward[curState][self.action[a]] + (self.gamma * tmp)
22         tmpValue.append(tmp)
23         if tmp > maxValue:
24             maxValue = tmp
25
26     for i in range(len(tmpValue)):
27         if math.fabs(maxValue - tmpValue[i]) < 1e-9:
28             newAction.append(i)
29     self.optimalPolicy[curState] = sorted(newAction)
30     for a in self.action.keys():
31         if self.action[a] in newAction:
32             self.policy[curState][self.action[a]] = 1.0 / len(newAction)
33         else:
34             self.policy[curState][self.action[a]] = 0.0
35     if oldAction != self.optimalPolicy[curState]:
36         policyStable = False
37     if policyStable:
38         break
39     numPass += 1

```

In the initialization phase, we initialize the policy as the random uniform policy as before. Otherwise, we will still get stuck. Accordingly, we need to set each states in `self.optimalPolicy` as north, east, south and west. The evaluation phase remains the same as iterative policy iteration.

In the improvement phase, we find the action that maximize the state value and update the random uniform policy until the policy get stable.

H3 3.2 Result

The algorithm converges faster than I estimated. The result is shown as:

```

Policy Iteration:
e      null    w      w      w      w
ne     n      nw     nw     nw     s
ne     n      nw     nw     es     s
ne     n      nw     es     es     s
ne     n      es     es     es     s
e      e      e      e      e      null

```

Fig.3 Policy Iteration

Respectively, the shortest distances from each state to the terminal states is:

Shortest Distance:					
1	0	1	2	3	4
2	1	2	3	4	4
3	2	3	4	4	3
4	3	4	4	3	2
5	4	4	3	2	1
5	4	3	2	1	0

Fig.4 Shortest Distance

We can easily verify the correctness of the algorithm.

H2 4. Value Iteration*

H3 4.1 Implementation

Thought not required, I implemented value iteration method to improve the policy. Value iteration is a method which update policy every iteration instead of update policy after evaluation. The algorithm takes the advantage of the Bellman optimality equation directly.

$$v_{k+1}(s) = \max_{a \in \mathcal{A}} (\mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^a v_k(s'))$$

And the implementation

```

1  def valueiteration(self):
2      #Initialize
3      .....
4      k = 0
5      #Loop
6      while True:
7          delta = 0
8          for curState in range(self.state):
9              oldValue = self.value[curState]
10             maxValue = -1.0e9
11             for a in self.action.keys():
12                 nextState = self.trans[curState][self.action[a]] #Only one element
13                 because here's GridWorld
14                 tmp = self.prob[curState][self.action[a]][nextState] *
15                 self.value[nextState] #prob must be 1.0
16                 tmp = self.reward[curState][self.action[a]] + (self.gamma * tmp)
17                 if tmp > maxValue:
18                     maxValue = tmp
19                 self.value[curState] = maxValue
20                 delta = max(delta, math.fabs(oldValue - self.value[curState]))
21             k += 1
22             if delta < self.threshold:
23                 break
24         for curState in range(self.state):
25             newAction = []

```

```

26     tmpValue = []
27     maxValue = -1.0e9
28     for a in self.action.keys():
29         nextState = self.trans[curState][self.action[a]] #Only one element
because here's GridWorld
30         tmp = self.prob[curState][self.action[a]][nextState] *
self.value[nextState] #prob must be 1.0
31         tmp = self.reward[curState][self.action[a]] + (self.gamma * tmp)
32         tmpValue.append(tmp)
33         if tmp > maxValue:
34             maxValue = tmp
35
36     for i in range(len(tmpValue)):
37         if math.fabs(maxValue - tmpValue[i]) < 1e-9:
38             newAction.append(i)
39     self.optimalPolicy[curState] = sorted(newAction)

```

The initialization phase is the same as policy iteration. In the while-loop, we update value function using Bellman optimality equation instead of Bellman equation. After that, we need to re-calculate the value function for each state and find the optimal policy.

H3 4.2 Result

We can get the result in only 7 iterations. The result is the same as it is in Policy Iteration section.

```

Value Iteration:
7
e      null    w      w      w      w
ne     n      nw     nw     nw     s
ne     n      nw     nw     es     s
ne     n      nw     es     es     s
ne     n      es     es     es     s
e      e      e      e      e      null

```

Fig.5 Value Iteration

H2 5. Summarize and Acknowledgement