

# H1 CS489 Assignment 4 Report

517030910214 Hongzhou Liu

## H2 0. Introduction

In this assignment, I implemented DQN and its improvement Double DQN in MountainCar environment. In MountainCar, there's a car on a one-dimension track. The goal is to drive up the mountain on the right. Due to the limitation of engine, the car has to drive back and forth to build up momentum in order to reach the goal. There are three kinds of actions:

- 0: push left
- 1: do nothing
- 2: push right

Each time the car chooses an action, it will gain a reward of -1. The state observed by agent from the environment consists of current horizontal position and velocity. The goal is at the position of 0.5. An episode ends when the car reaches the goal or the car takes 200 actions.




Fig.1 Mountain Car

Environment:

- Ubuntu 18.04 LTS
- Python 3.7.7

## H2 1. DQN

### H3 1.1 Implementation

#### H4 1.1.1 Algorithm

Deep Q-Network is a milestone of Reinforcement Learning. It introduced neural network into the original Q-Learning algorithm and largely strengthened the ability of Q-Learning algorithm. The two features of DQN are:

- Fixed target net
- Replay buffer

The DQN algorithm starts with policy network  $Q$  and target network  $\hat{Q}$ , which is the same. In each episode, the agent interact with the environment using the given state  $s_t$  and an action  $a_t$  based on  $Q$  epsilon greedily. The algorithm stores the obtained reward  $r_t$  and next state  $s_{t+1}$  together with  $s_t$  and  $a_t$  into the buffer as  $(s_t, a_t, r_t, s_{t+1})$ . Then, it samples a batch from the buffer and estimate the target as

$$y = r + \gamma \max_a \hat{Q}(s, a)$$

Then, we will optimize the mean square error between the estimation and Q value as

$$\min(y - Q(s, a))^2$$

The algorithm fixes the target network  $\hat{Q}$  and updates the policy network  $Q$ . And it will update  $\hat{Q}$  by setting  $\hat{Q} := Q$  every  $C$  steps.

#### H4 1.1.2 Network

We will utilize `torch.nn` to construct our neural network for DQN.

```
1 class DQN(nn.Module):
2     def __init__(self, inputSize, hiddenSize, outputSize):
3         super(DQN, self).__init__()
4         self.fc1 = nn.Linear(inputSize, hiddenSize)
5         self.fc2 = nn.Linear(hiddenSize, hiddenSize)
6         self.fc3 = nn.Linear(hiddenSize, outputSize)
7
8     def forward(self, x):
9         x = F.leaky_relu(self.fc1(x))
10        x = F.leaky_relu(self.fc2(x))
11        x = F.leaky_relu(self.fc3(x))
12        return x
```

Here, we construct a network with one hidden layer. And in this assignment, the size of the hidden layer is 256. We use leaky ReLU to avoid gradient vanishing. The input is state and the output is  $Q(s, a)$  for each action  $a$ .

#### H4 1.1.3 Replay Buffer

The replay buffer is like memory in our brain. It stores some recent experience in tuple  $(s_t, a_t, r_t, s_{t+1})$ . The size of the buffer is fixed, meaning it will forget as it interacts with the environment.

```
1 class ReplayBuffer:
2     def __init__(self, capacity):
3         self.capacity = capacity
4         self.buffer = []
5         self.position = 0 #for recursive array
6
7     def push(self, *args):
8         if len(self.buffer) < self.capacity:
9             self.buffer.append(None)
10            self.buffer[self.position] = Exp(*args)
11            self.position = (self.position + 1) % self.capacity
12
```

```

13     def sample(self, batchSize):
14         return random.sample(self.buffer, batchSize)

```

We can implement the buffer as a recursive array. If the buffer is full, the newly coming experiences will replace the old ones. Also, we can use `sample` method to sample a batch of experiences.

#### H4 1.1.4 Optimize Object Function

The kernel of Deep Q-Learning is to optimize the mean-squared error between  $Q$  and the estimation  $y$ . When implementing this part, some coding tricks are used.

```

1     def optimizeDQN(self, buf):
2         if len(buf) < self.batchSize:
3             print("Can't fetch enough exp!")
4             return
5         exps = buf.sample(self.batchSize)
6         batch = Exp(*zip(*exps)) # batch => Exp of batch
7         stateBatch = torch.cat(batch.state) # batchSize * stateSpace.shape[0]
8         actionBatch = torch.cat(batch.action) # batchSize * 1
9         rewardBatch = torch.cat(batch.reward) # batchSize * 1
10        nextStateBatch = torch.cat(batch.nextState) # batchSize *
stateSpace.shape[0]
11        doneBatch = torch.cat(batch.done) # batchSize * 1
12
13        Q = self.net(stateBatch).gather(1, actionBatch) # get Q(s_t, a)
14        targetQ = self.targetNet(nextStateBatch).max(1)[0].view(-1, 1) # max_a
Q'(s_{t+1}, a)
15        y = (targetQ * self.gamma) * doneBatch + rewardBatch
16        loss = F.mse_loss(Q, y)
17        estQMean = np.mean(targetQ.detach().cpu().numpy())
18
19        self.optimizer.zero_grad()
20        loss.backward()
21        self.optimizer.step()
22        return loss.item(), estQMean

```

After sampling a batch of experiences from buffer, we should separate the batches of  $s_t$ ,  $a_t$ ,  $r_t$  and  $s_{t+1}$  from the experience batch. Here, `zip()` function and `*` operator helps us reshape the batch and construct a new named tuple.

When calculating  $Q$ , `gather(1, actionBatch)` helps to find  $Q(s, a)$  according to `actionBatch` as indices. And `max(1)[0]` to get  $\max_a \hat{Q}(s, a)$ . Notice that  $y$  is slightly different when  $s_t$  is the last state of an episode. `doneBatch` consists of flags indicating if  $s_t$  is of such state. We multiply it on  $\gamma \max_a \hat{Q}(s, a)$  to make sure  $y = r_t$  if  $s_t$  is the last state.

#### H4 1.1.5 Learning

```

1     def DeepQLearning(self, env, buf, episodeNum, ifDecay=True):
2         totalN = 0
3         scores = []
4         steps = []
5         losses = []

```

```

6     Qmeans = []
7     self.targetNet.load_state_dict(self.net.state_dict())
8     for i in range(episodeNum):
9         score = 0.0
10        rewardDecay = 1.0
11        state = env.reset()
12        for t in count():
13            env.render()
14            action = self.getAction(torch.tensor(state.reshape(1, self.stateSize),
15device=device, dtype=torch.float))
16
17            nextState, reward, done, _ = env.step(action.item())
18
19            buf.push(torch.tensor(state.reshape(1, self.stateSize), device=device,
20dtype=torch.float),
21                    action,
22                    torch.tensor([[reward]], device=device, dtype=torch.float),
23                    torch.tensor(nextState.reshape(1, self.stateSize), device=device,
24dtype=torch.float),
25                    torch.tensor([[not done]], device=device, dtype=torch.long))
26            state = nextState
27            score += rewardDecay * reward
28            rewardDecay *= self.gamma
29
30            loss, mean = self.optimizeDQN(buf)
31            losses.append(loss)
32            Qmeans.append(mean)
33
34            if ifDecay:
35                self.epsilonDecay(totalN)
36                totalN += 1
37
38            if totalN % self.updateStride == 0:
39                self.targetNet.load_state_dict(self.net.state_dict())
40
41            if done or t + 1 >= env.max_episode_steps:
42                scores.append(score)
43                steps.append(t + 1)
44                print("Episode %d ended after %d timesteps with score %f,
45epsilon=%f" % (i + 1, t + 1, score, self.epsilon))
46                break

```

As seen, in the `for` loop, the DQN algorithm is properly implemented. And `self.getAction` implements epsilon greedy.

#### H4 1.1.6 Some Discussions

##### H5 Epsilon Decay

```

1     def epsilonDecay(self, N):
2         if (N >= 50000):
3             self.epsilon = self.EPS_END + (self.EPS_START - self.EPS_END) *
4                 math.exp(-1.0 * (N - 50000) / self.EPS_DECAY)

```

## H5 Customized Reward

```
1 def step(self, action):
2     state, reward, done, info = self.env.step(action)
3     if done:
4         reward = 1000
5     else:
6         reward = (self.height(state[0]) - 0.1) * 10 + 50 * state[1] * state[1]
7     return state, reward, done, info
8
9 def height(self, xs):
10    return np.sin(3 * xs) * .45 + .55
```

## H3 1.2 Result

## H2 2. Double DQN

### H3 2.1 Implementation

The double DQN is an improvement of natural DQN algorithm. It modified the method of estimating TD target and helps reduce overestimating of Q value greatly. The modified estimating of TD target is

$$y = r + \gamma \hat{Q}(s, \arg \max_a Q(s, a))$$

Now,

- If  $Q$  overestimate  $a$ , then  $\hat{Q}$  will give it a proper value
- If  $\hat{Q}$  overestimate  $a$ , then  $a$  will not be selected

Thus, we can avoid overestimation in this way.

To implement Double DQN, we can just modify the `optimizeDQN` function and change the part of calculating  $Q$  and  $y$  as

```
1 Q = self.net(stateBatch).gather(1, actionBatch) # get Q(s, a)
2 targetAction = self.net(nextStateBatch).max(1)[1].view(-1, 1) # get argmax
  Q'(s_{t+1}, a)
3 targetQ = self.targetNet(nextStateBatch).gather(1, targetAction)
4 y = (targetQ * self.gamma) * doneBatch + rewardBatch
5 loss = F.mse_loss(Q, y)
6 estQMean = np.mean(targetQ.detach().cpu().numpy())
```

As seen, now we use  $\arg \max_a \hat{Q}(s_{t+1}, a)$  as the action batch when estimating `targetQ`. Here, `max(1)[1]` is used to get the index of maximum value in a row of a tensor in two dimensions.

### H3 2.2 Result & Discussion